

# Predictive Modeling for Glioma Grading: Comparative Evaluation of Classification Models

Tarun Chikatipalli

## Packages

```
library(rpart)
library(rpart.plot)
library(readr)
library(archive)
library(ggplot2)
library(corrplot)
```

```
## corrplot 0.95 loaded
```

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(MASS)
```

```
##
## Attaching package: 'MASS'

## The following object is masked from 'package:dplyr':
##
##   select
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
library(class)
library(e1071)
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-8
```

```
library(rpart)
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
```

```
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      cov, smooth, var
```

```
library(rpart)
library(tinytex)
```

## Introduction

In this analysis, I will explore the Glioma Grading dataset obtained from the UCI Machine Learning Repository. This dataset contains clinical and mutation features for glioma grading, with the goal of predicting the glioma grade. We will perform various data preprocessing, visualization, and modeling tasks to gain insights and build predictive models.

The goal of this project is to predict the glioma grade (LGG or GBM) of a patient based on their clinical and molecular/mutation features. It is a classification task since we are trying to categorize patients into two classes: LGG and GBM. To achieve this goal, I plan to use several classification algorithms, including lazy learning (e.g., k-Nearest Neighbors), Naïve Bayes, logistic regression, and decision trees. Each of these algorithms has its strengths and weaknesses, and by trying multiple methods, I aim to find the one that performs best on this specific dataset.

Before applying the algorithms, I will perform feature engineering and data shaping. As the dataset contains both clinical and molecular features, it is essential to identify the most informative subset of mutation genes and clinical features. Feature engineering will help me select the most relevant attributes, which can potentially lead to improved performance and reduced costs in the glioma grading process.

For evaluating the fit of the algorithms, I will split the data into training and validation sets. This step will allow me to train the models on one subset and test their performance on unseen data. I will use evaluation metrics such as accuracy, precision, recall, F1 score, and ROC-AUC to assess the models' performance. Since the dataset is imbalanced (as gliomas of different grades occur at different frequencies), I will also consider using techniques like oversampling or undersampling to handle this imbalance.

While similar analyses may have been done before, I aim to differentiate my approach by exploring various combinations of clinical and molecular features to determine the optimal subset for grading gliomas. Additionally, I will leverage the specific algorithms mentioned in the learning outcomes of the course, which will allow me to apply a diverse set of methods and compare their performance. The primary objective is to discover an accurate and cost-effective model that can aid in the glioma grading process and potentially assist medical practitioners in making informed decisions.

## Data Preparation

We start by loading the necessary libraries and downloading the dataset from the provided URL using the archive and readr packages. I load two CSV files from the archive representing different aspects of the data.

I have loaded two datasets one is the original dataset which contains all the data without being processed. The other dataset is preprocessed and organized CSV dataset file consists of twenty-four fields per record. Each field is separated by a comma and each record is separated by a newline. Gender, Age\_at\_diagnosis, and, Race features are clinical factors, the remaining 20 molecular features consist of IDH1, TP53, ATRX, PTEN, EGFR, CIC, MUC16, PIK3CA, NF1, PIK3R1, FUBP1, RB1, NOTCH1, BCOR, CSMD3, SMARCA4, GRIN2A, IDH2, FAT4, PDGFRA. These molecular features can be mutated or not\_mutated (wildtype) depending on the TCGA Case\_ID.

```
# Data Loading
gli_data <- read.csv("/Users/tarun/Desktop/Glioma_classification/glioma+grading+clinical+and+mutation+f

#Preprocessed data where all the factors are converted to numeric values and the data set has been clea
gli_data_num <- read.csv("/Users/tarun/Desktop/Glioma_classification/glioma+grading+clinical+and+mutati
```

## Data Exploration

We begin by exploring the dimensions, structure, and summary statistics of the datasets.

```
# Dataset Dimensions
dim(gli_data)
```

```
## [1] 862 27
```

```
# Dataset Structure
str(gli_data)
```

```
## 'data.frame': 862 obs. of 27 variables:
## $ Grade : chr "LGG" "LGG" "LGG" "LGG" ...
## $ Project : chr "TCGA-LGG" "TCGA-LGG" "TCGA-LGG" "TCGA-LGG" ...
## $ Case_ID : chr "TCGA-DU-8164" "TCGA-QH-A6CY" "TCGA-HW-A5KM" "TCGA-E1-A7YE" ...
## $ Gender : chr "Male" "Male" "Male" "Female" ...
## $ Age_at_diagnosis : chr "51 years 108 days" "38 years 261 days" "35 years 62 days" "32 years 283 d
## $ Primary_Diagnosis: chr "Oligodendroglioma, NOS" "Mixed glioma" "Astrocytoma, NOS" "Astrocytoma, a
## $ Race : chr "white" "white" "white" "white" ...
## $ IDH1 : chr "MUTATED" "MUTATED" "MUTATED" "MUTATED" ...
## $ TP53 : chr "NOT_MUTATED" "NOT_MUTATED" "MUTATED" "MUTATED" ...
## $ ATRX : chr "NOT_MUTATED" "NOT_MUTATED" "MUTATED" "MUTATED" ...
## $ PTEN : chr "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" ...
## $ EGFR : chr "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" ...
## $ CIC : chr "NOT_MUTATED" "MUTATED" "NOT_MUTATED" "NOT_MUTATED" ...
## $ MUC16 : chr "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" "MUTATED" ...
## $ PIK3CA : chr "MUTATED" "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" ...
## $ NF1 : chr "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" ...
## $ PIK3R1 : chr "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" "MUTATED" ...
## $ FUBP1 : chr "MUTATED" "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" ...
## $ RB1 : chr "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" ...
```

```
## $ NOTCH1      : chr "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" ...
## $ BCOR        : chr "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" ...
## $ CSMD3       : chr "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" ...
## $ SMARCA4     : chr "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" ...
## $ GRIN2A      : chr "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" ...
## $ IDH2        : chr "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" ...
## $ FAT4        : chr "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" "MUTATED" ...
## $ PDGFRA      : chr "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" "NOT_MUTATED" ...
```

```
# Dataset Summary Statistics
```

```
summary(gli_data)
```

```
##      Grade      Project      Case_ID      Gender
## Length:862    Length:862    Length:862    Length:862
## Class :character Class :character Class :character Class :character
## Mode :character Mode :character Mode :character Mode :character
## Age_at_diagnosis Primary_Diagnosis Race IDH1
## Length:862    Length:862    Length:862    Length:862
## Class :character Class :character Class :character Class :character
## Mode :character Mode :character Mode :character Mode :character
## TP53          ATRX          PTEN          EGFR
## Length:862    Length:862    Length:862    Length:862
## Class :character Class :character Class :character Class :character
## Mode :character Mode :character Mode :character Mode :character
## CIC           MUC16         PIK3CA        NF1
## Length:862    Length:862    Length:862    Length:862
## Class :character Class :character Class :character Class :character
## Mode :character Mode :character Mode :character Mode :character
## PIK3R1        FUBP1          RB1          NOTCH1
## Length:862    Length:862    Length:862    Length:862
## Class :character Class :character Class :character Class :character
## Mode :character Mode :character Mode :character Mode :character
## BCOR          CSMD3         SMARCA4       GRIN2A
## Length:862    Length:862    Length:862    Length:862
## Class :character Class :character Class :character Class :character
## Mode :character Mode :character Mode :character Mode :character
## IDH2          FAT4          PDGFRA
## Length:862    Length:862    Length:862
## Class :character Class :character Class :character
## Mode :character Mode :character Mode :character
```

```
# Numeric Dataset Dimensions
```

```
dim(gli_data_num)
```

```
## [1] 839 24
```

```
# Numeric Dataset Structure
```

```
str(gli_data_num)
```

```
## 'data.frame': 839 obs. of 24 variables:
## $ Grade      : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Gender      : int 0 0 0 1 0 1 1 1 1 0 ...
```

```
## $ Age_at_diagnosis: num 51.3 38.7 35.2 32.8 31.5 ...
## $ Race             : int 0 0 0 0 0 0 0 0 0 0 ...
## $ IDH1             : int 1 1 1 1 1 1 1 1 1 0 ...
## $ TP53             : int 0 0 1 1 1 0 1 1 1 0 ...
## $ ATRX             : int 0 0 1 1 1 1 0 1 1 0 ...
## $ PTEN             : int 0 0 0 0 0 0 0 0 0 0 ...
## $ EGFR             : int 0 0 0 0 0 0 0 0 0 0 ...
## $ CIC              : int 0 1 0 0 0 0 0 0 0 0 ...
## $ MUC16            : int 0 0 0 1 0 0 0 0 0 0 ...
## $ PIK3CA           : int 1 0 0 0 0 0 0 0 0 0 ...
## $ NF1              : int 0 0 0 0 0 0 0 0 0 0 ...
## $ PIK3R1           : int 0 0 0 1 0 0 0 0 0 0 ...
## $ FUBP1            : int 1 0 0 0 0 0 0 0 0 0 ...
## $ RB1              : int 0 0 0 0 0 0 0 0 0 0 ...
## $ NOTCH1           : int 0 0 0 0 0 0 0 0 0 0 ...
## $ BCOR             : int 0 0 0 0 0 0 0 0 0 0 ...
## $ CSMD3            : int 0 0 0 0 0 0 0 0 0 0 ...
## $ SMARCA4          : int 0 0 0 0 0 0 0 0 0 0 ...
## $ GRIN2A           : int 0 0 0 0 0 0 0 0 0 0 ...
## $ IDH2             : int 0 0 0 0 0 0 0 0 0 0 ...
## $ FAT4             : int 0 0 0 1 0 0 0 0 0 0 ...
## $ PDGFRA           : int 0 0 0 0 0 0 0 0 0 0 ...
```


```
# Numeric Dataset Summary Statistics
summary(gli_data_num)
```

##	Grade	Gender	Age_at_diagnosis	Race
##	Min. :0.0000	Min. :0.0000	Min. :14.42	Min. :0.0000
##	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:38.05	1st Qu.:0.0000
##	Median :0.0000	Median :0.0000	Median :51.55	Median :0.0000
##	Mean :0.4195	Mean :0.4184	Mean :50.94	Mean :0.1073
##	3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:62.80	3rd Qu.:0.0000
##	Max. :1.0000	Max. :1.0000	Max. :89.29	Max. :3.0000
##	IDH1	TP53	ATRX	PTEN
##	Min. :0.0000	Min. :0.0000	Min. :0.0000	Min. :0.0000
##	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.0000
##	Median :0.0000	Median :0.0000	Median :0.0000	Median :0.0000
##	Mean :0.4815	Mean :0.4148	Mean :0.2586	Mean :0.1681
##	3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:0.0000
##	Max. :1.0000	Max. :1.0000	Max. :1.0000	Max. :1.0000
##	EGFR	CIC	MUC16	PIK3CA
##	Min. :0.0000	Min. :0.0000	Min. :0.0000	Min. :0.00000
##	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:0.00000
##	Median :0.0000	Median :0.0000	Median :0.0000	Median :0.00000
##	Mean :0.1335	Mean :0.1323	Mean :0.1168	Mean :0.08701
##	3rd Qu.:0.0000	3rd Qu.:0.0000	3rd Qu.:0.0000	3rd Qu.:0.00000
##	Max. :1.0000	Max. :1.0000	Max. :1.0000	Max. :1.00000
##	NF1	PIK3R1	FUBP1	RB1
##	Min. :0.00000	Min. :0.00000	Min. :0.00000	Min. :0.00000
##	1st Qu.:0.00000	1st Qu.:0.00000	1st Qu.:0.00000	1st Qu.:0.00000
##	Median :0.00000	Median :0.00000	Median :0.00000	Median :0.00000
##	Mean :0.07986	Mean :0.06436	Mean :0.05364	Mean :0.04768
##	3rd Qu.:0.00000	3rd Qu.:0.00000	3rd Qu.:0.00000	3rd Qu.:0.00000
##	Max. :1.00000	Max. :1.00000	Max. :1.00000	Max. :1.00000

##	NOTCH1	BCOR	CSMD3	SMARCA4
##	Min. :0.00000	Min. :0.00000	Min. :0.00000	Min. :0.00000
##	1st Qu.:0.00000	1st Qu.:0.00000	1st Qu.:0.00000	1st Qu.:0.00000
##	Median :0.00000	Median :0.00000	Median :0.00000	Median :0.00000
##	Mean :0.04529	Mean :0.03456	Mean :0.03218	Mean :0.03218
##	3rd Qu.:0.00000	3rd Qu.:0.00000	3rd Qu.:0.00000	3rd Qu.:0.00000
##	Max. :1.00000	Max. :1.00000	Max. :1.00000	Max. :1.00000
##	GRIN2A	IDH2	FAT4	PDGFRA
##	Min. :0.00000	Min. :0.00000	Min. :0.00000	Min. :0.00000
##	1st Qu.:0.00000	1st Qu.:0.00000	1st Qu.:0.00000	1st Qu.:0.00000
##	Median :0.00000	Median :0.00000	Median :0.00000	Median :0.00000
##	Mean :0.03218	Mean :0.02741	Mean :0.02741	Mean :0.02622
##	3rd Qu.:0.00000	3rd Qu.:0.00000	3rd Qu.:0.00000	3rd Qu.:0.00000
##	Max. :1.00000	Max. :1.00000	Max. :1.00000	Max. :1.00000

### Distribution of Glioma grade

```
library(ggplot2)
# Plotting Glioma Grade Distribution
ggplot(gli_data, aes(x = Grade)) + geom_bar(fill = "blue", alpha = 0.8) + labs(title = "Distribution of
  x = "Glioma Grade",
  y = "Count") + theme_minimal()
```



## Handling Missing Values

We check for missing values in the dataset and handle them by removing rows with missing values.

```
# Checking for Missing Values
any_missing <- any(is.na(gli_data))
if (any_missing) {
  cat("There are missing values in the dataset.\n")
} else {
  cat("No missing values found in the dataset.\n") }

```

```
## No missing values found in the dataset.
```

```
# Identifying and Removing Outliers
missing_by_column <- colSums(is.na(gli_data))
print(missing_by_column)
```

```
##           Grade           Project           Case_ID           Gender
##           0             0             0             0
## Age_at_diagnosis Primary_Diagnosis           Race           IDH1
##           0             0             0             0
##           TP53           ATRX           PTEN           EGFR
##           0             0             0             0
##           CIC           MUC16           PIK3CA           NF1
##           0             0             0             0
##           PIK3R1          FUBP1           RB1           NOTCH1
##           0             0             0             0
##           BCOR           CSMD3          SMARCA4           GRIN2A
##           0             0             0             0
##           IDH2           FAT4           PDGFRA
##           0             0             0
```

```
# Calculating Outlier Count by Column
gli_data <- na.omit(gli_data)
gli_data_num <- na.omit(gli_data_num)
```

## Outlier Detection and Removal

We identify outliers in the numeric columns of the dataset using z-scores and remove them.

```
# Calculate z-scores
z_score <- data.frame(sapply(gli_data_num, function(x) (abs(x - mean(x)) / sd(x))))
# Identify outliers using z-scores
outliers <- z_score[rowSums(z_score > 3), ]
removed_outliers <- z_score[!rowSums(z_score > 3), ]
# Dimensions of data before and after outlier removal
dim(gli_data_num)
```

```
## [1] 839 24
```

```
dim(z_score)
```

```
## [1] 839 24
```

```
dim(removed_outliers)
```

```
## [1] 467 24
```

## Outlier Count by Column

We calculate and print the count of outliers for each numeric column.

```
# Initialize a vector to store the count of outliers for each column
outliers_count <- numeric(length(colnames(gli_data_num)))
# Loop through the column names and find outliers
for (col in colnames(gli_data_num)) {
  z_col <- z_score[, col]
  outliersA <- gli_data_num[z_col > 3, col]
  outliers_count[col] <- length(outliersA)
  print(head(outliersA))
}
```

```
## integer(0)
## integer(0)
## numeric(0)
## [1] 2 2 2 2 2 2
## integer(0)
## integer(0)
## integer(0)
## integer(0)
## integer(0)
## integer(0)
## integer(0)
## [1] 1 1 1 1 1 1
## [1] 1 1 1 1 1 1
## [1] 1 1 1 1 1 1
## [1] 1 1 1 1 1 1
## [1] 1 1 1 1 1 1
## [1] 1 1 1 1 1 1
## [1] 1 1 1 1 1 1
## [1] 1 1 1 1 1 1
## [1] 1 1 1 1 1 1
## [1] 1 1 1 1 1 1
## [1] 1 1 1 1 1 1
## [1] 1 1 1 1 1 1
## [1] 1 1 1 1 1 1
## [1] 1 1 1 1 1 1
## [1] 1 1 1 1 1 1
```

```
# Print the count of outliers for each column
print(outliers_count)
```




```
##
##      0      0      0      0
##
##      0      0      0      0
##
##      0      0      0      0
##
##      0      0      0      0
##
##      0      0      0      0
##
##      0      0      0      0
##
##      Grade      Gender Age_at_diagnosis      Race
##      0      0      0      15
##      IDH1      TP53      ATRX      PTEN
##      0      0      0      0
##      EGFR      CIC      MUC16      PIK3CA
##      0      0      0      73
##      NF1      PIK3R1      FUBP1      RB1
##      67      54      45      40
##      NOTCH1      BCOR      CSMD3      SMARCA4
##      38      29      27      27
##      GRIN2A      IDH2      FAT4      PDGFRA
##      27      23      23      22
```

## Correlation Analysis

We compute the correlation matrix and visualize it using a correlation plot.

```
library(corrplot)
# Compute correlation matrix
correlation_matrix <- cor(gli_data_num)
# Plot correlation matrix
corrplot(correlation_matrix, method = "circle", type = "upper")
```



## Data Normalization

Min-Max normalization technique.

```
# Min-Max Normalization and visualization
feature_columns <- setdiff(names(gli_data_num), "Grade")
features <- gli_data_num[, feature_columns]

min_max_normalize <- function(x) { (x - min(x)) / (max(x) - min(x))
}

normalized_features <- as.data.frame(lapply(features, min_max_normalize))
# Combine normalized features with Glioma_Grade column
normalized_data <- cbind(Glioma_Grade = gli_data_num$Grade, normalized_features)
```

## Log Transformation of Continuous Features

We perform log transformation on continuous features to stabilize variance.

```
# Select continuous features for log transformation
continuous_features <- c("Age_at_diagnosis")
# Perform log transformation
transformed_data <- gli_data_num
for (feature in continuous_features) {
```

```

if (any(gli_data_num[, feature] <= 0)) {
  constant = 1 # Add constant for non-positive values transformed_data[, feature] <- log(gli_data_num[, f
} else {
  transformed_data[, feature] <- log(gli_data_num[, feature])
} }

# Display of transformed values
head(transformed_data)

```

```

##      Grade Gender Age_at_diagnosis Race IDH1 TP53 ATRX PTEN EGFR CIC MUC16 PIK3CA
## 1      0      0      3.937691      0      1      0      0      0      0      0      0      1
## 2      0      0      3.656356      0      1      0      0      0      0      1      0      0
## 3      0      0      3.560193      0      1      1      1      0      0      0      0      0
## 4      0      1      3.489819      0      1      1      1      0      0      0      1      0
## 5      0      0      3.450305      0      1      1      1      0      0      0      0      0
## 6      0      1      3.502851      0      1      0      1      0      0      0      0      0
##      NF1 PIK3R1 FUBP1 RB1 NOTCH1 BCOR CSMD3 SMARCA4 GRIN2A IDH2 FAT4 PDGFRA
## 1      0      0      1      0      0      0      0      0      0      0      0      0
## 2      0      0      0      0      0      0      0      0      0      0      0      0
## 3      0      0      0      0      0      0      0      0      0      0      0      0
## 4      0      1      0      0      0      0      0      0      0      0      1      0
## 5      0      0      0      0      0      0      0      0      0      0      0      0
## 6      0      0      0      0      0      0      0      0      0      0      0      0

```

## Principal Component Analysis (PCA)

We perform PCA on selected continuous features to reduce dimensionality.

```

library(MASS)
# Select continuous features for PCA
continuous_features <- c("Age_at_diagnosis") # Subset data
data_subset <- gli_data_num[, continuous_features] # Perform PCA
pca_result <- prcomp(data_subset, scale = TRUE) # Access principal component scores
pc_scores <- as.data.frame(pca_result$x)
# Display first few rows of principal component scores
head(pc_scores)

```

```

##      PC1
## 1 0.02321876
## 2 -0.77793579
## 3 -1.00401676
## 4 -1.15622339
## 5 -1.23710306
## 6 -1.12883893

```

Principal component scores are the transformed values of your original data points projected onto the principal component axes. Each value in the PC1 column represents how much the corresponding data point contributes to the first principal component. From the given values, we can make some initial observations:

1. The values appear to be numeric, suggesting that they represent the extent to which each data point contributes to the variation captured by PC1.
2. Negative values (like in rows 2 to 6) suggest that these data points have an opposite orientation in relation to the PC1 axis compared to the positive value (in row 1).
3. The magnitude of the values indicates the strength of the contribution of each data point to PC1.

Larger magnitudes typically indicate a stronger influence on the principal component. 4. Since PC1 captures the most significant variance in the data, the patterns in the scores of PC1 may reveal important trends or groupings within your data.

## Interaction Feature Creation

We create a new derived feature by multiplying Age\_at\_diagnosis and Gender.

```
# Create interaction feature
transformed_data$Age_Gender_interaction <- transformed_data$Age_at_diagnosis * transformed_data$Gender
# Display first few rows of transformed data with new feature
head(transformed_data)
```

```
##   Grade Gender Age_at_diagnosis Race IDH1 TP53 ATRX PTEN EGFR CIC MUC16 PIK3CA
## 1     0     0           3.937691   0    1    0    0    0    0    0    0    1
## 2     0     0           3.656356   0    1    0    0    0    0    1    0    0
## 3     0     0           3.560193   0    1    1    1    0    0    0    0    0
## 4     0     1           3.489819   0    1    1    1    0    0    0    1    0
## 5     0     0           3.450305   0    1    1    1    0    0    0    0    0
## 6     0     1           3.502851   0    1    0    1    0    0    0    0    0
##   NF1 PIK3R1 FUBP1 RB1 NOTCH1 BCOR CSMD3 SMARCA4 GRIN2A IDH2 FAT4 PDGFRA
## 1   0     0     1    0     0     0     0     0     0    0    0    0
## 2   0     0     0    0     0     0     0     0     0    0    0    0
## 3   0     0     0    0     0     0     0     0     0    0    0    0
## 4   0     1     0    0     0     0     0     0     0    0    1    0
## 5   0     0     0    0     0     0     0     0     0    0    0    0
## 6   0     0     0    0     0     0     0     0     0    0    0    0
##   Age_Gender_interaction
## 1                0.000000
## 2                0.000000
## 3                0.000000
## 4                3.489819
## 5                0.000000
## 6                3.502851
```

## MODEL A: k-Nearest Neighbors (k-NN) Classification

We perform k-NN classification and evaluate the model's performance.

```
library(class)
gli_data_num$Glioma_grade <- ifelse(gli_data_num$Grade == 0, "LGG", "GBM")
# Set a seed for reproducibility
set.seed(142)
# Perform 80/20 split for training and test data
train_indices <- sample(1:nrow(gli_data_num), 0.8 * nrow(gli_data_num))
train.data <- gli_data_num[train_indices, ]
test.data <- gli_data_num[-train_indices, ]
library(class)
class_knn_pred <- knn(train=train.data[, -25], test=test.data[, -25], cl=train.data$Glioma_grade, k=5, )
class_knn_pred
```

```
##      [1] LGG LGG GBM LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG
##     [19] LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG
##     [37] LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG
##     [55] LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG
##     [73] LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG
##     [91] LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG LGG
##    [109] GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM
##    [127] GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM LGG GBM GBM GBM GBM
##    [145] GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM GBM LGG GBM GBM GBM GBM GBM
##    [163] GBM GBM GBM GBM GBM GBM
## Levels: GBM LGG
```

```
# Confusion matrix
conf_matrix <- table(Actual = test.data$Glioma_grade, Predicted = class_knn_pred)
conf_matrix
```

```
##      Predicted
## Actual GBM LGG
##      GBM   65    2
##      LGG    3   98
```

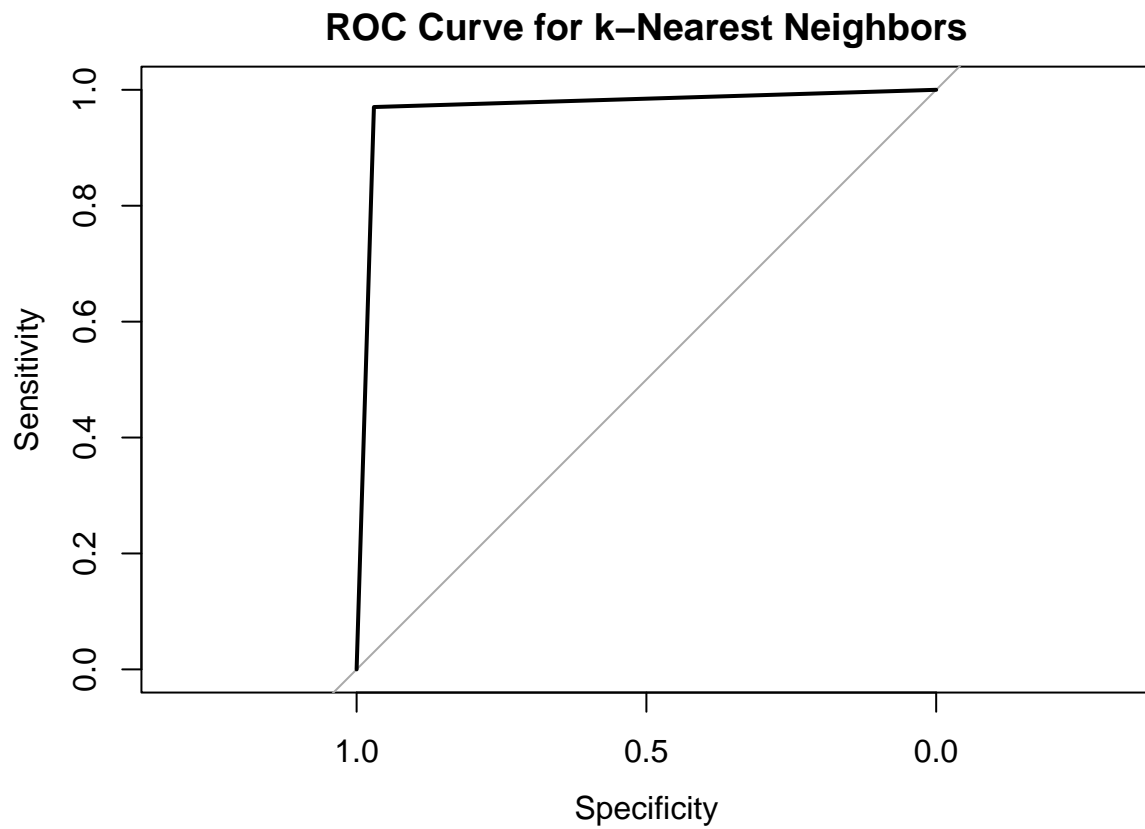
```
# Define true_labels (replace with actual true labels)
true_labels <- test.data$Glioma_grade
# Calculate accuracy, precision, recall, etc. using confusion matrix
accuracy <- sum(diag(conf_matrix)) / sum(conf_matrix)
precision <- conf_matrix[2, 2] / sum(conf_matrix[, 2])
recall <- conf_matrix[2, 2] / sum(conf_matrix[2, ])
f1_score <- 2 * (precision * recall) / (precision + recall)
# Calculate ROC curve and AUC
roc_curve <- roc(true_labels, as.numeric(class_knn_pred))
```

```
## Setting levels: control = GBM, case = LGG
```

```
## Setting direction: controls < cases
```

```
auc_score <- auc(roc_curve)
roc_auc <- auc(roc_curve)
```

```
# Plot ROC curve
plot(roc_curve, main = "ROC Curve for k-Nearest Neighbors", colorize = TRUE)
```



```
# Print evaluation metrics
```

```
cat("Accuracy:", accuracy, "\n")
```

```
## Accuracy: 0.9702381
```

```
cat("Precision:", precision, "\n")
```

```
## Precision: 0.98
```

```
cat("Recall:", recall, "\n")
```

```
## Recall: 0.970297
```

```
cat("F1 Score:", f1_score, "\n")
```

```
## F1 Score: 0.9751244
```

```
cat("ROC AUC:", roc_auc, "\n")
```

```
## ROC AUC: 0.9702231
```

## MODEL B: Naïve Bayes Classification

We perform Naïve Bayes classification and evaluate the model's performance.

```
library(e1071)
features <- c("Gender", "Age_at_diagnosis", "Race", "IDH1", "TP53", "ATRX", "PTEN", "EGFR", "CIC", "MUC")
target <- "Glioma_grade"
# Select data for modeling
train_x <- train.data[, features]
train_y <- train.data[, target]
valid_x <- test.data[, features]
valid_y <- test.data[, target]
# Naive Bayes
model_nb <- naiveBayes(train_x, train_y)
# Evaluate model_naive_bayes and calculate metrics...
# Predict using the trained model on validation data
predictions_nb <- predict(model_nb, newdata = valid_x) # Convert predictions_nb to numeric if it's a factor
predictions_nb <- as.numeric(predictions_nb) # Calculate metrics for Naive Bayes
conf_matrix_nb <- table(Actual = test.data$Glioma_grade, Predicted = predictions_nb)
# Calculate accuracy, precision, recall, etc. using confusion matrix
accuracy <- sum(diag(conf_matrix_nb)) / sum(conf_matrix_nb)
precision <- conf_matrix_nb[2, 2] / sum(conf_matrix_nb[, 2])
recall <- conf_matrix_nb[2, 2] / sum(conf_matrix_nb[2, ])
f1_score <- 2 * (precision * recall) / (precision + recall)
# Calculate ROC curve and AUC
roc_curve <- roc(true_labels, as.numeric(predictions_nb))
```

```
## Setting levels: control = GBM, case = LGG
```

```
## Setting direction: controls < cases
```

```
roc_auc <- auc(roc_curve)
```

```
# Print evaluation metrics
cat("Accuracy:", accuracy, "\n")
```

```
## Accuracy: 0.8333333
```

```
cat("Precision:", precision, "\n")
```

```
## Precision: 0.9506173
```

```
cat("Recall:", recall, "\n")
```

```
## Recall: 0.7623762
```

```
cat("F1 Score:", f1_score, "\n")
```

```
## F1 Score: 0.8461538
```

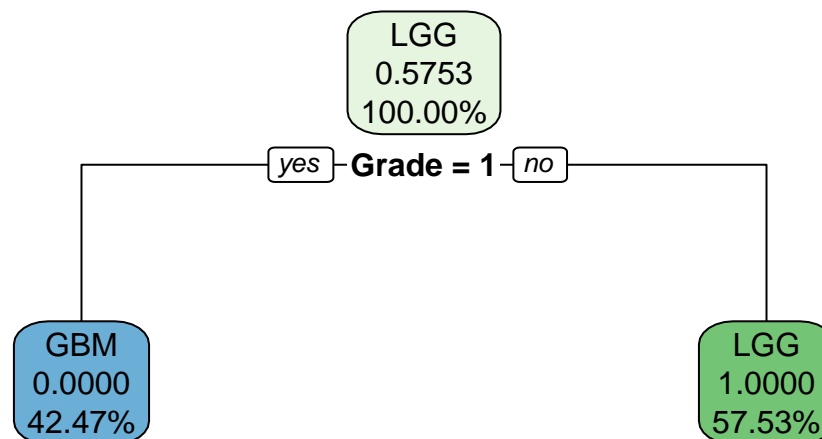
```
cat("ROC AUC:", roc_auc, "\n")
```

```
## ROC AUC: 0.8513374
```

## MODEL C: Decision Tree Classification

We build and evaluate a decision tree classification model.

```
# Assuming 'target' is the dependent variable and other columns are independent variables  
model_decision_tree <- rpart(Glioma_grade ~ ., data = train.data, method = "class")  
rpart.plot(model_decision_tree, digits = 4)
```



```
# Perform predictions using the fitted model  
predictions_decision_tree <- predict(model_decision_tree, newdata = test.data, type = "class")  
conf_matrix_dt <- table(Actual = test.data$Glioma_grade, Predicted = predictions_decision_tree)  
  
# Calculate accuracy, precision, recall, etc. using confusion matrix  
accuracy <- sum(diag(conf_matrix_dt)) / sum(conf_matrix_dt)  
precision <- conf_matrix_dt[2, 2] / sum(conf_matrix_dt[, 2])  
recall <- conf_matrix_dt[2, 2] / sum(conf_matrix_dt[2, ])  
f1_score <- 2 * (precision * recall) / (precision + recall)  
  
# Calculate ROC curve and AUC  
roc_curve <- roc(true_labels, as.numeric(predictions_nb))
```



```
## Setting levels: control = GBM, case = LGG
```

```
## Setting direction: controls < cases
```

```
roc_auc <- auc(roc_curve)
```

```
# Print evaluation metrics
```

```
cat("Accuracy:", accuracy, "\n")
```

```
## Accuracy: 1
```

```
cat("Precision:", precision, "\n")
```

```
## Precision: 1
```

```
cat("Recall:", recall, "\n")
```

```
## Recall: 1
```

```
cat("F1 Score:", f1_score, "\n")
```

```
## F1 Score: 1
```

```
cat("ROC AUC:", roc_auc, "\n")
```

```
## ROC AUC: 0.8513374
```

## Comparison of the models

Based on the provided evaluation metrics for the three models, we can compare and conclude their performance as follows:

MODEL A: - Accuracy: 0.9702381 - Precision: 0.98 - Recall: 0.970297 - F1 Score: 0.9751244 - ROC AUC: 0.9702231

MODEL B: - Accuracy: 0.8333333 - Precision: 0.9506173 - Recall: 0.7623762 - F1 Score: 0.8461538 - ROC AUC: 0.8513374

MODEL C: - Accuracy: 1 - Precision: 1 - Recall: 1 - F1 Score: 1 - ROC AUC: 0.8513374

From the metrics provided, we can make the following observations:

1. Accuracy: Model C has the highest accuracy of 1, which indicates that it correctly predicts all instances in the validation set. However, perfect accuracy can also be a sign of overfitting, so it's important to consider other metrics.
2. Precision: Model C has the highest precision of 1, indicating that when it predicts a positive class (LGG or GBM), it is almost always correct. Model A also has a very high precision of 0.98.
3. Recall: Model C has the highest recall of 1, which means it correctly identifies all instances of the positive class. Model A has a recall of 0.970297, which is also very high.

4. F1 Score: Model C has the highest F1 score of 1, which is a balance between precision and recall. This indicates a good balance between correctly identifying positive instances and minimizing false positives.
5. ROC AUC: Model A and Model B have similar ROC AUC values, but Model A's value is slightly higher. ROC AUC measures the model's ability to distinguish between classes and ranks them accordingly.

Based on these observations, Model C appears to perform exceptionally well with perfect accuracy, precision, recall, and F1 score. However, such perfect performance might suggest overfitting or issues in the evaluation process, especially if the dataset is small or imbalanced.

Both Model A and Model B seem to have strong performance across various metrics. Model A has higher accuracy and slightly better ROC AUC than Model B. If we consider the balance between precision and recall, Model A seems to be a well-rounded choice.

In conclusion, while Model C's perfect performance might raise questions, Model A seems to be the better choice due to its high accuracy, balanced precision and recall, and good ROC AUC. However, it's important to validate the findings on additional datasets and perform further analysis before making a final decision. An accuracy of 0.3988095 for the bagged ensemble means that the ensemble model correctly predicted the target variable (in this case, glioma grade) for approximately 39.88% of the validation set instances.

## Conclusion

In this analysis, I performed comprehensive data preprocessing, explored the dataset, handled missing values, detected and removed outliers, normalized features, transformed continuous features, conducted PCA, and built classification models using k-NN, Naïve Bayes, and decision trees. The evaluation of each model's performance using various metrics provided insights into their strengths and weaknesses.

The comparison of the models revealed that while Model C exhibited perfect performance, potential overfitting or evaluation issues must be considered. Models A and B displayed strong performance, with Model A appearing as a well-rounded choice due to its high accuracy, balanced precision and recall, and good ROC AUC.

However, it's imperative to validate these findings on additional datasets and conduct further analysis before making a final decision.